




Retrieval Confidence vs Answer Quality (RAG rows only)




ROUGE-1 by Model and Embedding Type
(temp=0.0, k=5 — dashed = no-RAG baseline)



LLM Judge Score by Model and Embedding Type (temp=0.0, k=5)




RAG vs No-RAG Baseline (temp=0.0, k=5 — Mistral has no no-RAG condition)




Fine-tuning Delta by Question Type
(temp=0.0, k=5)


ROUGE-1 Delta (Fine-tuned - Base)



Effect of Temperature on ROUGE-1
(Lower temp → more focused answers)




Effect of Retrieval Depth k on ROUGE-1 (temp=0.0)




Mean Latency by Model and Embedding Type

* Mistral = local GPU; GPT/Claude = API calls



Per-Company ROUGE-1 Heatmap
(Fine-tuned embeddings, temp=0.0, k=5)



Numeric Accuracy on Numerical Questions (temp=0.0, k=5)

